


Intro and Markov Models

- value of information: $EV(x \text{ is known}) - EV(x \text{ is unknown})$ [known in this case means an arc from x to a decision variable]
- value of control: $EV(x \text{ is a decision variable}) - EV(x \text{ is a random variable})$
- **limitation of decision networks:**
 - represents only a fixed # of decisions
- **advantage of markov models:**
 - network can extend indefinitely

Markov Models



A stationary Markov Chain : for all $t > 0$

- $P(S_{t+1} | S_0, \dots, S_t) = P(S_{t+1} | S_t)$ and Markov assumption
- $P(S_{t+1} | S_t)$ is the same stationary

So we only need to specify?

iclicker.

A. $P(S_{t+1} | S_t)$ and $P(S_0)$

B. $P(S_0)$

Value Iteration

Value iteration works by producing successive approximations of the optimal value function.

$$\forall s: V^{(k+1)}(s) = R(s) + \gamma \max_a \sum_{s'} P(s' | s, a) V^{(k)}(s')$$

What is the complexity of each iteration?

A. $O(|A|^2|S|)$

B. $O(|A||S|^2)$

C. $O(|A|^2|S|^2)$

Iteration 1				
3	0	0	0.76	$V^{(1)}(3,3) = -0.04 + 1 * \max \begin{cases} 0.8V^{(0)}(3,3) + 0.1V^{(0)}(2,3) + 0.1V^{(0)}(4,3) & \text{UP} \\ 0.8V^{(0)}(2,3) + 0.1V^{(0)}(3,3) + 0.1V^{(0)}(3,2) & \text{LEFT} \\ 0.8V^{(0)}(3,2) + 0.1V^{(0)}(2,3) + 0.1V^{(0)}(4,3) & \text{DOWN} \\ 0.8V^{(0)}(4,3) + 0.1V^{(0)}(3,3) + 0.1V^{(0)}(3,2) & \text{RIGHT} \end{cases}$
2	0		0	
1	-0.04	0		
	1	2	3	4

Markov Decision Processes

- set of states S , actions A , transition probabilities and reward function

Decision Processes

Often an agent needs to go beyond a fixed set of decisions – Examples?

- Would like to have an ongoing decision process

Infinite horizon problems: process does not stop

Robot surviving on planet, Monitoring Nuc. Plant, ...

Indefinite horizon problem: the agent does not know when the process may stop

Finite horizon: the process must end at a give time N

- MDPs are fully observable
- Agent always knows which state its in
- Uncertainty lies in its actions

FILTERING

- Policy only depends on current state
- In general (Filtering): compute the posterior distribution over the current state given all evidence to date

$$P(X_t | e_{0:t})$$

POMDP belief states:

$$b^i(s') = \alpha P(e^i | s') \sum_s P(s' | a, s) b(s)$$

- how to find optimal policy?
- turn into MDP and apply V.I.
- also approx. methods ie DDN using look ahead

Look Ahead Search for Optimal Policy

General Idea:

- Expand the decision process for n steps into the future, that is
 - "Try" all actions at every decision point
 - Assume receiving all possible observations at observation points
- Result: tree of depth $2n+1$ where
 - every branch represents one of the possible sequences of n actions and n observations available to the agent, and the corresponding belief states
 - The leaf at the end of each branch corresponds to the belief state reachable via that sequence of actions and observations – use filtering/belief-update to compute it
- "Back Up" the utility values of the leaf nodes along their corresponding branches, combining it with the rewards along that path
- Pick the branch with the highest expected value

Reinforcement Learning

- transition and reward model are unknown, can't exploit relation to adjacent states
 - this is addressed through temporal difference
 - takes an average through time for Q-values
- A belief state of an agent encodes what an agent has remembered
- Discounting means that more recent rewards are more valuable than rewards far in the future.

Q-Learning

- Search based applications' policies are evaluated as a whole, and cannot take into account local good/bad actions
- n states, m actions = m^n
- Q-Learning learns after every action for what to do in a given state

- Learning is linear instead of exponential with # of states

$$Q[s, a] \leftarrow Q[s, a] + \alpha((r + \gamma \max_{a'} Q[s', a']) - Q[s, a])$$

- For the approximation in Q-Learning to work, try each action an unbounded # of times

- Unlikely transitions will be observed much less frequently than likely ones

- Any strategy should be greedy in the limit of infinite exploration (GLIE)

- Choose the predicted best action in the limit, and try an unbounded # of times

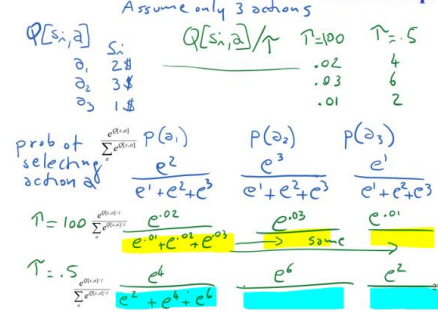
E-greedy

- choose random action with probability E , choose best action with probability $1-E$
- reduce E overtime to predict best action in GLIE

Soft-max

- High τ = exploration
- Low τ = exploitation

Soft-Max example



On and Off-Policy learning

- If an agent is not deployed it should do random all the time and Q-Learning
 - Deploy once Q values have converged
- If an agent is deployed use one of the explore/exploit strategies and do SARSA
- Q-Learning is off-policy; it focuses on the optimal policy
 - On-policy learning addresses problem of high risk exploration and can revise policy

Hoeffding's inequality

You can tolerate an error greater than 0.1 only in 5% of your cases

Set $\epsilon = 0.1$, $\delta = 0.05$

Equation (1) gives you $n > 184$

$$n > \frac{-\ln \frac{\delta}{2}}{2\epsilon^2} \quad (1)$$

Filtering

$$P(X_1 | e_1) = \alpha P(e_1 | X_1) * \sum_{x_0} P(X_1 | x_0) P(x_0)$$

$$= \alpha (0.8 \cdot 0.3) * (0.5 \cdot (0.7 \cdot 0.3) + 0.5 \cdot (0.4 \cdot 0.6))$$

Backwards

$$P(e_3 | X_2) = \sum_{x_3} P(e_3 | x_3) P(x_3 | X_2)$$

$$= (0.8 * 1 * (0.7 \cdot 0.4)) + (0.3 * 1 * (0.3 \cdot 0.6))$$

Smoothing

$$P(X_k | e_{0:k}) = P(X_k | e_{0:k}, e_{k+1:t}) \quad \text{dividing up the evidence}$$

$$= \alpha P(X_k | e_{0:k}) P(e_{k+1:t} | X_k, e_{0:k}) \quad \text{using Bayes Rule}$$

$$= \alpha P(X_k | e_{0:k}) P(e_{k+1:t} | X_k) \quad \text{By Markov assumption on evidence}$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

$$b_{t+1:t} = P(e_{t+1:t} | X_t) = P(\text{unspecified} | X_t) = 1$$

- Limitations of exact algorithms
- HMM has very large # of states
- Algorithms do not scale up
 - Use **Approximate inference**
 - **Getting N samples is faster than computing the right answer (via filtering)**

- **Particle filtering:** run all N samples together through the network one slice at a time

1. Generate a population on N initial-state samples by sampling from initial state distribution $P(x_0)$
2. Propagate each sample for x_t forward by sampling the next state value x_{t+1} based on $P(x_{t+1} | x_t)$
3. Weight each sample by the likelihood it assigns to the evidence
4. Resample the population so that the probability that each sample is selected is proportional to its weight
5. Start this cycle again with this new sample

HMM Application

Bioinformatics: Gene Finding

• States: coding / non-coding region $\times \times \checkmark \checkmark \checkmark \checkmark$

• Observations: DNA Sequences $\rightarrow ATCGGAA$

For these problems the critical inference is:

find the most likely sequence of states given a sequence of observations

Graphical Models

- Variable elimination algorithm for Bnets and Markov networks are the same!
- A CRF models $P(y_1, \dots, y_k | x_1, \dots, x_k)$
- special case of Markov Networks where x_i 's are all observed
- it is an undirected graphical model whose nodes = $X \cup Y$

- logistic regression is a naïve markov model

- number of param w_i is linear instead of exponential the # of parents

- natural model

- natural model

Markov Networks Applications (1): Computer Vision

Called **Markov Random Fields**

• Stereo Reconstruction

• Image Segmentation

• Object recognition

Typically **pairwise MRF**

• Each vars correspond to a **pixel** (or **superpixel**)

• Edges (factors) correspond to interactions between adjacent pixels in the image

• E.g., in segmentation: from generically penalize discontinuities, to road under car

Propositional logic and Resolution

Validity and satisfiability

A sentence is valid if it is true in all interpretations

e.g., $\text{True}, A \vee \neg A, A \Rightarrow A, (A \wedge (A \Rightarrow B)) \Rightarrow B$

Validity is connected to inference via the **Deduction Theorem:**

$KB \models \alpha$ if and only if $(KB \Rightarrow \alpha)$ is valid

A sentence is **satisfiable** if it is true in some interpretation

e.g., $A \vee B, C$

A sentence is unsatisfiable if it is true in no interpretations

e.g., $A \wedge \neg A$

Satisfiability is connected to inference via the following:

$KB \models \alpha$ if and only if $(KB \wedge \neg \alpha)$ is unsatisfiable

i.e., prove α by **reductio ad absurdum**

- Resolution algorithm tries to prove $KB \models \alpha$

- Convert to CNF, apply resolution
- Until two clauses resolves into an empty clause
- > contradiction, therefore unsatisfiable
- No new clauses can be added; no contradiction

Walksat

- Start from randomly generated implementation
- Pick a proposition/atom to flip by
 - 1) Randomly
 - 2) To minimize # of unsatisfied clauses
- problems get hard as #clauses/#symbols ↑ (~4.3)
- resolution procedure can be generalized to **FOL**
- every formula can be rewritten in a logically equivalent CNF form

- additional rewriting rules for quantifiers
- variables need to be unified (like datalog)

Wordnet – database containing lexical relations

- To compute similarities between two senses in WordNet, use:

- the distance between the two concepts in the underlying hierarchies/graphs
- the glosses of the concepts
- **lesk** can be disadvantageous because most glosses are short and don't have overlaps
- Path-length sim based on is-a/hypernyms hierarchies

$$\text{sim}_{\text{path}}(c_1, c_2) = 1 / \text{pathlen}(c_1, c_2)$$

- comparing two words would be semantically ambiguous
- comparing two synsets we are confident it's close in distance/overlap


Relations among words and their meanings (*paradigmatic*)

Internal structure of individual words (*syntagmatic*)

- FrameNet** – database containing frames and syntactic/semantic argument structures
- depth vs breadth vs density (breadth and density = understanding)

Concept Distance: info content

- Similarity should be proportional to the information that the two concepts share... what is that?



$$\text{sim}_{\text{resnik}}(c_1, c_2) = -\log P(LCS(c_1, c_2))$$

$$\text{dist}_{\text{JC}}(c_1, c_2) = 2 \times \log P(LCS(c_1, c_2)) - (\log P(c_1) + \log P(c_2))$$

Treebanks

- DEF. corpora in which each sentence has been paired with a parse tree
- These are generally created
 - Parse collection with parser
 - human annotators revise each parse
- Requires detailed annotation guidelines
 - POS tagset
 - Grammar
 - instructions for how to deal with particular grammatical constructions.

WS Distributional Methods (2)

- More informative values (referred to as weights or measure of association in the literature)

- Point-wise Mutual Information

$$\text{assoc}_{\text{PMI}}(w, w_i) = \log_2 \frac{P(w, w_i)}{P(w)P(w_i)}$$

if independently unrelated

$$\text{assoc}_{\text{t-test}}(w, w_i) = \frac{P(w, w_i) - P(w)P(w_i)}{\sqrt{P(w)P(w_i)}} \sqrt{\frac{N}{N-1}}$$

CPSC 412, Lecture 24

Context Free Grammar

CFGs

- Define a **Formal Language** (un/grammatical sentences)

- **Generative Formalism**

- Generate strings in the language
- Reject strings not in the language
- Impose structures (trees) on strings in the language

Top-Down vs. Bottom-Up

Top-Down

- Only searches for trees that can be answers
- But suggests trees that are not consistent with the words

Bottom-Up

- Only forms trees consistent with the words
- Suggest trees that make no sense globally

- CFGs cover most syntactic structures in English

- vanilla PCFGs assume independence of non-terminal expansions - **Invalid assumption!**

- structural/lexical dependencies
 - syntactic subject of a sentence as a pronoun
 - Jim likes dogs. But he lost one as a child.
 - **solution:** split non-terminals
 - add lexical dependencies to the scheme
 - infiltrate influence of particular words into the prob. of the rules

CKY Algorithm

Definitions

- $w_1 \dots w_n$ an input string composed of n words
- w_{ij} a string of words from word i to word j
- $\mu(i, j, A)$: a table entry holds the maximum probability for a constituent with non-terminal A spanning words $w_i \dots w_j$

Markov Logic

Prob. Rel. Models vs. Markov Logic

PRM


- Relational Skeleton
 - Dependency Graph
 - Parameters (CPT)
- \Rightarrow BNET

ML

- weighted logical formulas
 - set of constants
- \Rightarrow MARKOV LOGIC NETWORK

Markov Logic: Intuition(1)

- A logical KB is a set of **hard constraints** on the set of possible worlds



- Let's make them **soft constraints:** When a world violates a formula, the world becomes less probable, not impossible

if ω is True $P(\omega)$ decreases
 if ω is False $P(\omega)$ increases

Markov Logic: Definition

- A **Markov Logic Network (MLN)** is

- a set of pairs (F, w) where
 - F is a formula in first-order logic
 - w is a real number

- Together with a set C of constants,

- It defines a **Markov network** with

- One **binary node** for each **grounding** of each predicate in the MLN
- One **feature/factor** for each **grounding** of each formula F in the MLN, with the corresponding weight w

Grounding: substituting vars with constants

- prob. of possible world:

$$P(x) = \frac{1}{Z} \exp \left(\sum_i w_i f_i(x) \right)$$

Weight of Feature i Feature i

First order logic (with some mild assumptions) is a special Markov Logics obtained when

- all the weight are equal
- and tend to infinity

Inference in MLN

- MLN acts as a template for a Markov Network
- We can always answer prob. queries using standard Markov network inference methods on the instantiated network
- **However**, due to the size and complexity of the resulting network, this is often infeasible.
- Instead, we combine **probabilistic methods** with ideas from **logical inference**, including **satisfiability** and **resolution**.
- This leads to efficient methods that take full advantage of the logical structure.

Finding truth assignment that maximizes sum of weights of satisfied formulas:

$$\arg \max_{pw} \sum_i w_i n_i(pw)$$

-> **this is just a weighted MaxSAT problem!**

- instead of minimizing, maximize # of satisfied formulas

Entity Resolution (ML application)

- determining which observations correspond to the same real-world objects
- can perform **collective** entity resolution, where resolving one pair of entities can help resolve other relevant entities

Statistical Parsing

Representation: Summary

- For each rule of the form $A \rightarrow B C$:
Formula of the form $B(i, j) \wedge C(j, k) \Rightarrow A(i, k)$
E.g.: $\text{NP}(i, j) \wedge \text{VP}(j, k) \Rightarrow \text{S}(i, k)$
- For each rule of the form $A \rightarrow a$:
Formula of the form $\text{Token}(a, i) \Rightarrow A(i, i+1)$
E.g.: $\text{Token}(\text{"pizza"}, i) \Rightarrow \text{N}(i, i+1)$
- For each nonterminal: state that exactly one production holds (solve problem 1)
- Mutual exclusion rules for each part of speech pair (solve problem 2)

PRMs


- The representation of PRMs is a direct mapping from that of relational databases

Motivation for PRMs

- Most real-world data are stored in relational DBMS
- Combine advantages of relational logic & Bayesian networks:
 - natural domain modeling: objects, properties, relations;
 - generalization over a variety of situations;
 - compact, natural probability models.
- Integrate uncertainty with relational model:
 - properties of domain entities can depend on properties of related entities;
 - uncertainty over relational structure of domain.

Limitations of Bayesian Networks

A Bayesian networks (BNs) represents a pre-specified set of attributes/variables whose relationship to each other is fixed in advance.



How PRMs extend BNs?

1. PRMs conceptually extend BNs to allow the specification of a probability model for classes of objects rather than a fixed set of simple attributes
 2. PRMs also allow properties of an entity to depend probabilistically on properties of other related entities
- Each class has a set of attributes + primary keys
 - A set of **reference slots** (corresponds to foreign keys) ie. *Course.instructor* -> *Professor.name*
 - An instance has values for each attribute and reference slot (fixed vs probabilistic)
 - ie. name, id vs ranking, grade
 - Skeleton is a partial specification of an instance (doesn't have probabilistic att.)
 - Completion I of the skeleton is a possible world, specifies values for probabilistic attributes
 - A student is registered in many courses - student is a field in registration, Registered_in() is an **inverse slot** of a reference slot

Slot-chain: *Student.Registered_in.Course.Instructor*

- PRM contains a CPD for each attribute
- The parameters in all these CBDs comprise θ
- Problem: There are variable # of parents
 - Solution: use aggregation
 - Example: a pro gamer's ELO can depend on average KDA
- Given a skeleton, apply local cond. probabilities to define a Joint prob. dist. over all completions of the skeleton
- Disallow uncertainty over the relational structure of the model because of our skeleton structure x
 - Parameter sharing/CPT reuse also found in temporal models (stationary assumption!)
 - We must ensure our probabilistic dependencies are acyclic
 - This is guaranteed by some prior knowledge about the domain
 - Male chromosome from a person